Recalibrating Visual GEO: Anchor-Decoupled Evaluation and a Transfer Probe

Anonymous Author(s)
Submission under triple-blind review
ICDM 2026

***Abstract*—Visual generative engine optimization edits a product image so it ranks higher under a CLIP-style retriever, and headline results report large cosine-similarity gains at high SSIM, arguing that visual GEO is a serious retrieval-marketplace threat. We show much of this is an artifact of anchor-coupled measurement: when the evaluation query is itself derived from the optimization anchor, any such cosine gain is structurally inflated, shrinking to a bounded, cohort-specific effect under controlled decoupling. We make two contributions to honest threat measurement. First, an anchor-decoupled evaluation protocol — pinned by a single anchor⊥query invariant ($v_{\mathrm{anchor}}(x_p)$ depends only on x_p), a pre-flight retriever-discrimination gate, a per-cohort decomposition over the four ESCI human-graded relevance classes, and a compact reference verifier (released as HONESTEVAL) — that re-measures, rather than amplifies, the white-box threat of our per-image attack CoGEO against three baselines. Second, AE-CoGEO, a query-free, anchor-decoupled cross-encoder transfer *probe* for visual GEO. Its ensemble objective is standard; under the query-free/anchor-decoupled/retrieval-rank constraint triple it yields a calibrated transfer ceiling with a monotone source-count scaling law (ρ=1.0). On OpenAI ViT-L/14, our CoGEO is the strongest of the four attacks and the unique top-1 on five of six retrievers, its white-box advantage concentrated in the harder cohorts and the aggregate mean (+19.86, paired Wilcoxon $p \leq 1.1 \times 10^{-4}$, 560-pair intersection); the per-pair effect is right-skewed (median ≈1), a cohort-driven, tail-concentrated mean rather than a uniform shift. It trades off monotonically with stealth, and its perturbation transfers far below the white-box ceiling. Properly measured, visual GEO is a cohort-specific, white-box-bounded, transfer-limited, and detectable (AUROC ≥ 0.99) operator, not a uniform invisible threat.**

***Index Terms*—visual retrieval, CLIP, adversarial perturbation, generative engine optimization, Amazon ESCI, multimodal evaluation**

I. INTRODUCTION

Visual generative engine optimization is the seller-side analogue of search-engine optimization for retrieval systems that match natural-language queries to product images through a shared multimodal embedding [1], [2]. Imperceptible perturbations of a product image can substantially change its rank under a CLIP-style retriever [3]–[5], and headline results report large positive Δ in cosine similarity together with high SSIM [6], arguing that visual GEO is a real threat to retrieval marketplaces.

The conventional visual-GEO protocol is unsound in a reproducible way: the eval-time query is derived from the same anchor that drove the optimization. When $v_{\mathrm{anchor}}(x_p)$ is the CLIP encoding of a description of x_p and similarity is measured against the same caption, any positive Δ is a tautology of gradient ascent: optimization vigor, not retrieval impact under a real shopper query. The eval-time query must therefore be *decoupled* from the optimization anchor.

This paper makes two contributions. **First**, we deliver *HonestEval*, a reusable anchor-decoupled protocol that recalibrates visual GEO into an actionable, defender-relevant threat characterization rather than a deflated headline number. Figure 1 sketches the three-stage pipeline (seller anchor → four-family optimization → held-out ESCI scoring on white-box and transfer axes), with an Anchor-Decoupled Query Firewall keeping the query out of the gradient. The protocol pins three commitments: an *anchor⊥query invariant* making the anchor a function of the product image alone, checked by a compact ∼60-line verifier; an *ESCI-grounded benchmark* [7] of 491 images balanced across the four classes (E/S/C/I) with 560 paired triples, replicated across six CLIP backbones; and a *four-way comparison* of our proposed attack CoGEO against PGD-bare [8], AdvCLIP [4], and Co-Attack [5] at $\varepsilon = 16/255$, with per-cohort decomposition and a paired Wilcoxon matrix [9]. **Second**, to reach the deployment-realistic transfer axis the white-box protocol cannot, we contribute *AE-CoGEO* (§III-D), the first query-free, anchor-decoupled cross-encoder transfer probe for visual GEO, optimizing one perturbation against a source-encoder consensus to generalize to an unseen deployed encoder.

Three claims survive the protocol. **(C1)** *Visual GEO's threat scope is bounded and cohort-specific.* Under the anchor-decoupled protocol, our CoGEO is the strongest of the four compared CLIP-targeted attacks on the white-box axis, but as a cohort-driven aggregate-mean effect, not a per-pair median win: +19.86 mean rank-lift on the four-way intersection with paired Wilcoxon $p \leq 1.1{\times}10^{-4}$ (carried by the harder cohorts; the per-pair median lift is ≈ 1 and the advantage over PGD-bare goes non-significant at 1,430 pairs, one-sided p=0.78; §VI-E, VI-I) — not a state-of-the-art attack claim. **(C2)** Aggregate dominance hides per-cohort heterogeneity: CoGEO leads by large margins on Complement (+36.8) and Irrelevant

Code, the ESCI manifest, the no-model admissibility verifier, the reference rank-lift scorer, reference implementations of both attacks, reference result tables, and the full reproduction configuration (sampling, seeds, hyper-parameters, hardware) are released anonymously as HONESTEVAL: https://anonymous.4open.science/r/honestvgeo-F663.

(+25.4); on the Exact cohort the mean rank-lift gap reverses (+3.9 for CoGEO vs. +7.1 for PGD-bare), but matched-pair Wilcoxon and sign tests do not reject the null, so neither method has a significant per-pair Exact advantage despite the aggregate-mean gap. Co-Attack additionally turns negative on Complement, so single-number visual-GEO benchmarks conceal the per-cohort operating regime. **(C3)** Visual stealth and retrieval lift trade off across families: SSIM means rank in broadly reverse order from rank-lift (AdvCLIP 0.78 > PGD-bare/Co-Attack 0.76 > CoGEO 0.66), and CoGEO's SSIM ≥ 0.98 holds at 800×800 but not the 224×224 regime industrial CLIP retrieval consumes [1].

The boundary against prior work is twofold. On the evaluation side, any visual-GEO author can run the verifier of §III-A before reporting rank-lift, and the multi-method, multi-backbone replication [10]–[12] establishes the per-cohort pattern as an attack-family property, not any single embedding's. On the attack side, AE-CoGEO is positioned as a diagnostic transfer *probe*, not a competing attack: its constraint triple (query-free + anchor-decoupled + retrieval-rank), not its standard ensemble objective, is the increment (Table I, §II). We assume a seller-side attacker with read access to one or more *public* CLIP encoders and unlimited per-image gradient compute, lifting an upload's rank against held-out queries on the same encoder (white-box) or an unseen deployed one (transfer). The transfer measurement (§VI-G) and a training-free detector characterize visual GEO as a cohort-specific, white-box-bounded, transfer-limited operator.

II. Related Work

Generative engine optimization: GEO was introduced for text-only generative search engines [2], where small content edits change which sources a system cites, sharing roots with SEO and adversarial information retrieval. Visual GEO extends it to product retrieval ranked by a multimodal embedding rather than a generative LM [1], [13]. Our CoGEO attack combines a seller-side (product-title) anchor decoupled from the eval query, a Sobel mask, a differentiable JPEG layer [14], and an MI-FGSM loop [15]; we additionally ablate VLM-captioner anchors [16], [17]; it is one of four compared methods.

CLIP-targeted adversarial attacks: Universal adversarial patches [4], [18] on a CLIP image encoder shift retrieval across inputs, and a complementary family attacks both branches with alternating updates [5]; both are query-agnostic, matching the seller-side regime where uploads outpace per-image budgets. The per-image branch descends from PGD [8]/FGSM [3] with momentum [15], and recent work maps contrastive-pretraining vulnerability [19], [20] and robust-CLIP defenses [21]. What is missing is evaluation under held-out human relevance: prior work scores against queries derived from the optimization signal itself, never a multi-cohort benchmark with a published anchor⊥query invariant.

Transfer-based attacks and cross-encoder consensus: A long line improves black-box transferability of per-image attacks through input-diversity [22] and scale-invariant [23] transforms, momentum [15], and surrogate ensembling [24]. These were developed for classification and still presume a target label or query/caption at attack time [5]. AE-CoGEO (§III-D) is, to our knowledge, the first to bring ensemble-consensus transfer to visual GEO under the query-free, anchor-decoupled constraint, scoring rank-lift on held-out encoders absent from the optimization set (§VI-G). Table I positions it among these families.

TABLE I
THE FOUR EVALUATION AXES VISUAL GEO IMPOSES, AGAINST WHICH PRIOR ATTACK EVALUATION WAS NEVER MEASURED. **QF**: QUERY/CAPTION-FREE; **HE**: SCORED UNDER HELD-OUT HUMAN RELEVANCE (ANCHOR⊥QUERY); **XE**: TRANSFERS TO AN UNSEEN ENCODER VIA SOURCE-SET CONSENSUS; **VG**: OPTIMIZES RETRIEVAL RANK, NOT CLASSIFICATION. ONLY OUR ANCHOR-DECOUPLED PROTOCOL CAN *score* THE **HE** AXIS (HENCE EVERY PRIOR ROW IS BLANK THERE); AE-COGEO IS MEASURABLE ON ALL FOUR.

Attack	QF	HE	XE	VG
PGD / FGSM [3], [8]	✓	×	×	×
MI/DI/SI-FGSM [15], [22], [23]	✓	×	×	×
Ensemble transfer [24]	✓	×	✓	×
AdvCLIP patch [4]	✓	×	×	✓
Co-Attack [5]	×	×	×	✓
CoGEO (ours)	✓	✓	×	✓
AE-CoGEO (ours)	✓	✓	✓	✓

Held-out relevance evaluation: The Amazon ESCI dataset [7] provides (q, x_p, ℓ) triples with four-class human labels (Exact, Substitute, Complement, Irrelevant), and to our knowledge has not previously been used to evaluate adversarial image perturbations; we add a Food-101 [25] second-domain replication. Because ESCI queries come from real shoppers and labels are independent of any retriever's geometry, observed rank-lift measures something other than the gradient that produced it, placing the protocol in the held-out tradition of BEIR [26] and MTEB [27] that rejects self-loop similarity [21], [28].

Stealth and replication: SSIM [6] is the canonical perceptual-quality measure, but visual-GEO stealth claims are often reported at the optimization resolution (800×800) rather than the 224×224 input the retriever consumes — a discrepancy not previously characterized across attack families on one benchmark. We report the four-method stealth–lift Pareto at native input size and replicate across six CLIP backbones [1], [10]–[12].

III. Anchor-Decoupled Evaluation Protocol

A. The anchor⊥query invariant

Let x_p denote a product image, q a shopper query in natural language, E_I and E_T the image and text encoders of a CLIP-style retriever [1], [10], [13], and $v_{\text{anchor}}(x_p) \in \mathbb{R}^d$ a target text embedding used to drive a perturbation δ such that $E_I(x_p+\delta)$ moves toward $v_{\text{anchor}}(x_p)$.

We require that, for every (q, x_p) pair used in evaluation,

$$v_{\text{anchor}}(x_p) \text{ is a function of } x_p \text{ alone, not of } q. \quad (1)$$

Concretely, v_{anchor} may be the CLIP text encoding of the product category, of a VLM description of x_p (varied across

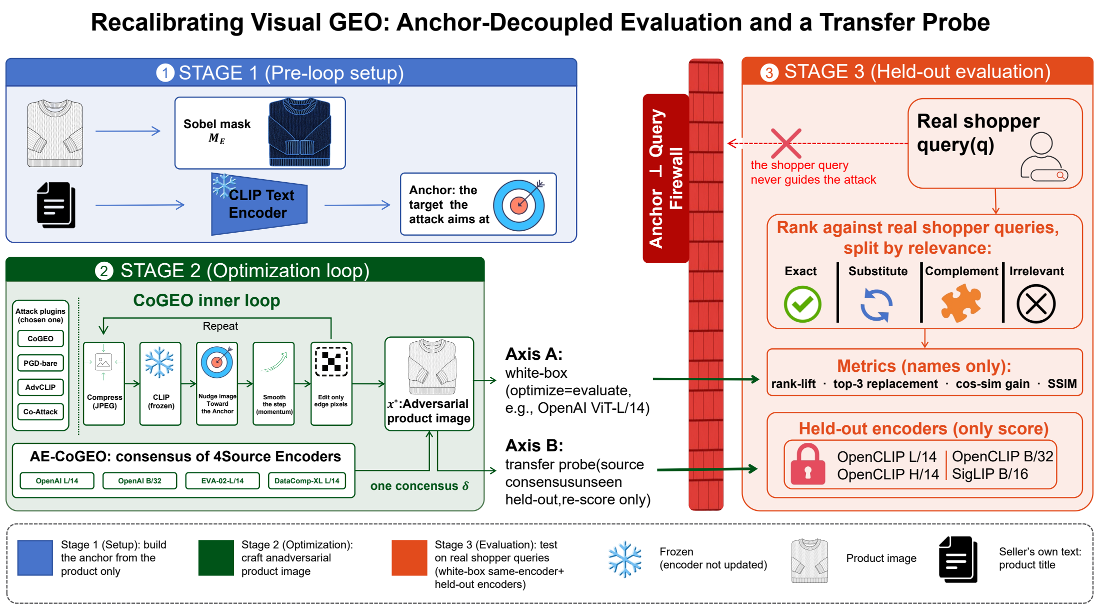

Fig. 1. Anchor-decoupled evaluation protocol with a cross-encoder transfer probe, instantiated for CoGEO. **Stage 1** precomputes the seller anchor $v_{\text{anchor}}{=}E_T(t_p^\star)$ from the product title and a Sobel-edge mask M_E (top-40% gradient pixels); v_{anchor} depends on the product image alone, never on a shopper query. **Stage 2** runs the CoGEO inner loop—JPEG-compress, encode with the frozen CLIP tower, nudge toward the anchor, smooth with momentum, edit only masked edge pixels—for $K{=}200$ iterations at $\varepsilon{=}16/255$; baselines PGD-bare, AdvCLIP and Co-Attack drop or replace individual boxes. Two orthogonal axes, not competing attacks: *Axis A* optimizes against a single white-box encoder; *Axis B* (AE-CoGEO) against a four-source consensus loss $L_{\text{ens}}{=}-\frac{1}{|S|}\sum_j \cos(E_I^{(j)}(\tilde{x}), E_T^{(j)}(t_p^\star))$. **Stage 3** scores the frozen image against unseen ESCI shopper queries across the four cohorts (E/S/C/I) and re-scores Axis B—no re-attack—on four held-out encoders. The shopper query never crosses the Anchor-Decoupled Query Firewall. Snowflakes mark frozen modules.

BLIP, BLIP-2, LLaVA-1.5, Qwen2-VL, and Qwen2.5-VL in the released artifact), or of the product title; in all cases it depends only on the seller-side description. The query q enters only at evaluation time, when we measure rank-lift, top-3 replacement, and held-out cosine similarity gains under q.

This invariant is the operational definition of an honest visual-GEO evaluation: when the eval query is itself derived from v_{anchor}, the similarity gain is structurally inflated relative to held-out retrieval impact. We enforce (1) at the code level across all four families with a compact reference verifier ($\sim$60 lines, released artifact) that statically rejects any query-like attack parameter and dynamically checks that the perturbation is byte-identical across swapped eval queries under a fixed seed.

B. Pre-flight monotonicity gate

Before any adversarial perturbation, the chosen retriever must demonstrate that human-graded relevance labels are discriminable in its embedding space. Writing $\overline{\cos}_\ell = \text{mean}_{(q,x_p)\in\ell}\big[\cos(E_I(x_p),\, E_T(q))\big]$ for the mean no-perturbation cosine over a cohort ℓ, we require that on a balanced sample of (q, x_p, ℓ) triples spanning the four ESCI relevance classes Exact (E), Substitute (S), Complement (C), Irrelevant (I), the gate

$$\overline{\cos}_{\ell=E} - \overline{\cos}_{\ell=I} \;\geq\; 0.05 \tag{2}$$

holds. A retriever that fails (2) cannot serve as an evaluation substrate: no adversarial pressure is measurable against an embedding that cannot already separate Exact-match from Irrelevant pairs. We read 0.05 as a qualitative admissibility floor, not a calibrated statistic: a conservative margin that the canonical CoGEO retriever (ViT-L/14, gap 0.0533) clears and a markedly weaker substrate (ViT-B/32, 0.0405) does not, with both satisfying $E > S > C > I$ monotonicity. The evaluation substrate (ViT-L/14) is fixed independently as the canonical retriever, so no downstream result depends on the exact floor value. On failure, the protocol mandates a backbone swap before optimization.

C. Four-way attack pass

Within the invariant, four white-box attack families are evaluated at a matched perturbation budget $\varepsilon = 16/255$, the cross-family normalization point at which all four families separate cleanly (not a claim that it is the only deployment-relevant budget; the budget-dependence is examined in §VI-D):

CoGEO: Per-image projected gradient descent [3], [8] on x_p with momentum (MI-FGSM [15]), Sobel-edge mask restricting perturbation to high-gradient regions, environment simulation through a differentiable JPEG layer [14] together with input-diversity [22] and scale-invariant [23] transformations, and product-title v_{anchor}; CoGEO is the attack we propose and the primary subject of the held-out evaluation. Adapting it to the anchor-decoupled protocol fixes three design choices, each settled by an ablation in the released artifact: a hard top-40% Sobel-gradient mask (vs. a soft min-max-normalized mask floored at 0.1), differentiable JPEG at Q=75 (vs. Q=85), and translation-invariant smoothing omitted, as it gave no significant lift under the anchor⊥query setting.

Formally, let $x_p^{(k)}$ denote the listing image at iteration k (with $x_p^{(0)}$=x_p). The Sobel edge mask

$$M_E \;=\; \mathbf{1}\big[\,|\nabla_{\text{Sobel}}\, x_p| \;\geq\; q_{0.6}\big(|\nabla_{\text{Sobel}}\, x_p|\big)\,\big] \tag{3}$$

selects the top-40% gradient pixels (threshold $q_{0.6}$, the 60-th percentile of edge magnitude) and is computed once per image. At each iteration k, CoGEO applies a differentiable JPEG layer at quality Q=75 to make the attack compression-aware, embeds the result with the frozen CLIP image encoder, and forms the anchor-alignment loss

$$L^{(k)} \;=\; -\,\cos\!\big(\,E_I(\text{diffJPEG}_{Q=75}(x_p^{(k)})),\; v_{\text{anchor}}\,\big). \tag{4}$$

The signed gradient is smoothed by a momentum buffer with $\mu = 1.0$ following the original MI-FGSM formulation [15]:

$$g^{(k+1)} \;=\; \mu\, g^{(k)} \;+\; \frac{\nabla_x\, L^{(k)}}{\|\nabla_x\, L^{(k)}\|_1}, \quad \mu = 1.0, \tag{5}$$

The masked signed-momentum step (step size $\alpha = 4/255$) is then projected onto the ℓ_∞ ε-ball $\mathcal{B}_\infty(x_p, \varepsilon)$=$\{x : \|x - x_p\|_\infty \leq \varepsilon\}$:

$$x_p^{(k+1)} \;=\; \Pi_{\mathcal{B}_\infty(x_p,\,\varepsilon)}\Big(x_p^{(k)} - \alpha \cdot \text{sign}\big(M_E \odot g^{(k+1)}\big)\Big). \tag{6}$$

After K=200 iterations, $x_p^\star$ is the adversarial image evaluated against the held-out query in §III-E.

PGD-bare: The same product-title v_{anchor} as CoGEO but without the Sobel mask, environment simulation, or momentum, recovering the classical signed-step PGD baseline of [8] and isolating the anchor's contribution from the engineering layered on top.

AdvCLIP: A query-agnostic universal adversarial patch [18], [29] trained on the CLIP image encoder [4] over 30 epochs with budget $\varepsilon = 16/255$, applied to each x_p at evaluation. AdvCLIP requires no per-image gradient pass at deployment time.

Co-Attack: A cross-modal alternating attack [5] that interleaves an L_∞ image perturbation ($\varepsilon = 16/255$) and an L_2 text-anchor perturbation ($\varepsilon_v = 0.1$) over 200 iterations. The text branch operates on the seller-side anchor only, never on q, so the invariant is preserved.

All four pipelines take only x_p as optimization input; the held-out ESCI query never enters any optimization gradient, verified at the code level.

The CoGEO inner loop is fully specified by Eqs. (3)–(6); the full evaluation-pass pseudocode is in the released artifact.

D. AE-CoGEO: a cross-encoder-consensus transfer probe

The four-way comparison above is a *white-box* measurement: the encoder that drives the perturbation also scores it. This bounds one axis — how far an attacker can move the ranking of a retriever it fully controls — but leaves unmeasured the realistic deployment axis, where the platform retriever is an *unseen* encoder the attacker never optimized against. We introduce **AE-CoGEO** (Anchor-decoupled Ensemble CoGEO) to measure this second axis.

AE-CoGEO keeps both protocol commitments of CoGEO unchanged — the perturbation is *query-free* and *anchor-decoupled*, so the held-out shopper query never enters optimization ((1)) — and adds a single mechanism: instead of aligning to one encoder, it optimizes one perturbation against the *consensus* of a source set of encoders $\mathcal{S} = \{(E_I^{(j)}, E_T^{(j)})\}_{j=1}^m$ ($m = |\mathcal{S}| = 4$ source encoders). Writing $t_p^\star = \textsc{SellerAnchor}(x_p)$ for the seller-side anchor text and $\tilde{x} = \text{diffJPEG}_Q(x_p + \delta)$, the anchor-alignment loss is averaged over the source set,

$$L_{\text{ens}}(\delta) = -\frac{1}{m}\sum_{j=1}^{m} \cos\!\big(E_I^{(j)}(\tilde{x}),\, E_T^{(j)}(t_p^\star)\big), \tag{7}$$

and the masked MI-FGSM update of (5)–(6) is applied to $\nabla_\delta L_{\text{ens}}$ without change, so δ raises anchor alignment *simultaneously* across architectures rather than overfitting one encoder's geometry. Ensemble optimization is a standard transferability route [15], [22], [24]; AE-CoGEO's increment is combining it diagnostically with the query-free, anchor-decoupled constraint visual GEO imposes.

Source and held-out encoders are strictly disjoint: the source set is $\mathcal{S}$ = {OpenAI ViT-L/14, OpenAI ViT-B/32, EVA-02-L/14, DataComp-XL ViT-L/14} and the held-out set is {OpenCLIP-L/14, -H/14, -B/32 (all LAION-2B), SigLIP ViT-B/16}, with no shared checkpoint and no shared pretraining corpus. Every other setting — $\varepsilon = 16/255$, $K = 200$, Sobel mask, differentiable-JPEG and input-diversity environment simulation, and product-title anchor — is identical to the single-encoder CoGEO arm, so any change in held-out behavior is attributable to the consensus mechanism alone.

CoGEO and AE-CoGEO thus instrument two orthogonal axes of the same threat — the white-box and transfer ceilings — with AE-CoGEO the calibrated transfer probe; §VI-G reports the measurement.

E. Per-cohort decomposition and significance

For each adversarial image $x_p^* = x_p + \delta$ we compute four metrics against the held-out query q: *rank-lift* $\text{rank}(x_p, q) - \text{rank}(x_p^*, q)$, the positions x_p^* rises for q; *top-3 replacement rate*, the fraction of pairs where x_p^* enters the top 3 while x_p did not; *held-out* $\Delta_{\text{clip\_sim}} = \cos(E_I(x_p^*), E_T(q)) - \cos(E_I(x_p), E_T(q))$, the direct CLIP-similarity gain under q; and *SSIM* between x_p and x_p^* at the retriever's native resolution, the stealth at the regime industrial retrieval consumes. Every metric is then *decomposed* per cohort over the four ESCI relevance classes. Aggregate means across all pairs are reported alongside per-cohort means, never as a substitute. We further compute a paired Wilcoxon signed-rank test for each of the $\binom{4}{2} = 6$ method pairs on the same (q, x_p) index, so each pairwise comparison is verifiable independently of any cohort selection.

IV. EXPERIMENTAL SETUP

We instantiate the protocol of Section III on the Amazon ESCI shopper-query benchmark and run all four attack families against the same set of product images.

A. Benchmark and sampling

We sample (q, x_p, ℓ) triples from the Amazon ESCI dataset [7] (US locale), with $\ell \in \{E, S, C, I\}$ a four-level human-graded relevance label authored by real shoppers and independent of any retriever embedding. A label-balanced sample (random seed 42) yields 491 unique product images and 560 paired triples — 125 Exact, 125 Substitute, 185 Complement, 125 Irrelevant — each resized to 224×224 [1] (CDN resolution in the released artifact). We additionally replicate the protocol on Food-101 [25] as a ranking-internal-consistency check (Fig. 2f); because its cohorts are CLIP-text-similarity quartiles over the 101 class names, they are retriever-induced and not promoted to ESCI-equivalent human-graded relevance evidence.

B. Retriever and budgets

The retriever is OpenAI CLIP ViT-L/14 [1] used identically as both the optimization target and the evaluation embedder, fixed a priori as the canonical CoGEO retriever; it also clears the pre-flight gate of (2) ($\text{gap}_{E-I} = 0.0533 > 0.05$), whereas a weaker ViT-B/32 (0.0405) would not (full backbone sensitivity table in the released artifact). Cross-backbone replication uses five additional CLIP-style retrievers from the OpenCLIP toolbox [10] that vary pretraining corpus and architecture: OpenCLIP-L/14 (LAION-2B), EVA-02-L/14 [11], SigLIP ViT-B/16 [12], OpenCLIP ViT-B/32, and OpenCLIP ViT-H/14. Anchor-source robustness is measured separately in the released artifact using five VLM captioner anchors — BLIP [16], BLIP-2 [30], LLaVA-1.5 [31], Qwen2-VL [17], Qwen2.5-VL [32] — over the same six backbones (30 CoGEO cells).

For each held-out query q_i, rank-lift is computed against a fixed per-query candidate pool of size C=491 (the union of all manifest ASIN images), making it a within-pool diagnostic over the 491-image manifest rather than an absolute marketplace-scale rank: *rank* is the position of $x^{\star}_{p,i}$ in the descending $\cos(E_I(\cdot), E_T(q_i))$ sort. All four families operate at matched $\varepsilon = 16/255$ in L_∞, with $\varepsilon \in \{4, 8\}/255$ sweeps in Fig. 2b. CoGEO and PGD-bare run $K = 200$ momentum signed-step PGD iterations [8], [15]; CoGEO adds a differentiable JPEG layer [14] and input-diversity transforms [22], [23]. AdvCLIP trains its universal patch [4], [18], [29] over the same manifest used at evaluation, a regime favorable to AdvCLIP, so CoGEO's dominance is conservative. Co-Attack runs 200 alternating iterations at image $\varepsilon = 16/255$, text $\varepsilon_v = 0.1$ [5], its text branch perturbing a fixed seller-side anchor (*not* the held-out query). Anchor sources are held constant within method (product-title for CoGEO/PGD-bare, image-only for AdvCLIP; full hyper-parameters in the released artifact).

TABLE II
FOUR-WAY COMPARISON ON THE 560-PAIR JOINT SAMPLE AT $\varepsilon = 16/255$ WITH CLIP VIT-L/14. RANK-LIFT IS THE MEAN POSITIONS AN IMAGE RISES AGAINST ITS HELD-OUT ESCI QUERY; SSIM IS AT THE 224×224 INPUT; p-VALUES ARE ONE-SIDED PAIRED WILCOXON VERSUS COGEO; I→TOP-3 GAIN IS AGAINST THE UNPERTURBED BASELINE.

Method	Rank-lift	p vs CoGEO	SSIM	I→top-3 gain
CoGEO	**19.86**	—	0.66	**+4.1 pp**
PGD-bare	12.37	1.1×10^{-4}	0.76	+2.9 pp
AdvCLIP	6.45	3.5×10^{-30}	**0.78**	+1.6 pp
Co-Attack	5.77	1.3×10^{-8}	0.76	+2.9 pp

C. Significance protocol

For each comparison we report the paired Wilcoxon signed-rank test [9] on rank-lift, paired bootstrap 95% CIs (1000 resamples), and the held-out $\Delta_{\text{clip\_sim}}$ mean; SSIM [6] is reported at the native 224×224 input. Directional tests are one-sided; the six pairwise p-values are reported uncorrected, but the headline comparisons survive Holm–Bonferroni at the six-test count (the weakest, $p = 1.1 \times 10^{-4}$, clears α and indeed $\alpha/6$). The full 4×4 matrix is in the released artifact (§VI-B).

V. MAIN RESULTS

Table II reports the headline four-way comparison on the full 560-pair intersection. These results measure the *white-box axis*, where optimization and scoring encoders coincide; the deployment-realistic *transfer axis* (AE-CoGEO, §III-D) is measured separately in §VI-G. Three findings emerge.

Aggregate-mean dominance is significant: CoGEO's $+19.86$ mean rank-lift is $1.6\times$ the next-best baseline, with every pairwise paired Wilcoxon p-value well below 0.05 (to 3.5×10^{-30} versus AdvCLIP) on this 560-pair set — though the CoGEO−PGD-bare margin specifically attenuates to non-significance at the 1,430-pair scale-up (p=0.78, §VI-I). AdvCLIP's universal patch [4] moves the ranking least — a single δ against a fixed text concept does not adapt to per-image content [18] — and Co-Attack falls statistically below PGD-bare ($p = 2.4 \times 10^{-6}$), its text-anchor branch not compensating for the loss of CoGEO's mask and momentum [15], [22], [23].

Visual stealth and retrieval lift trade off: The four SSIM values [6] rank in the opposite order from rank-lift (Table II): at 800×800 CoGEO's SSIM is ≥ 0.98, but at the 224×224 regime the retriever consumes all four fall to low SSIM (CoGEO 0.66, stealthiest AdvCLIP 0.78; Section VI-C). The right-most column of Table II reports the metric a defender ultimately cares about — pushing Irrelevant products into the top-3 a shopper sees, against a 9.1% unperturbed baseline — where CoGEO translates its largest aggregate lift into top-3 placement.

VI. ANALYSIS

The aggregate result of §V hides a structurally heterogeneous picture, decomposed below along cohort, stealth, budget, backbone, transfer, and anchor robustness, then consolidated under a bounded-under-stress synthesis.

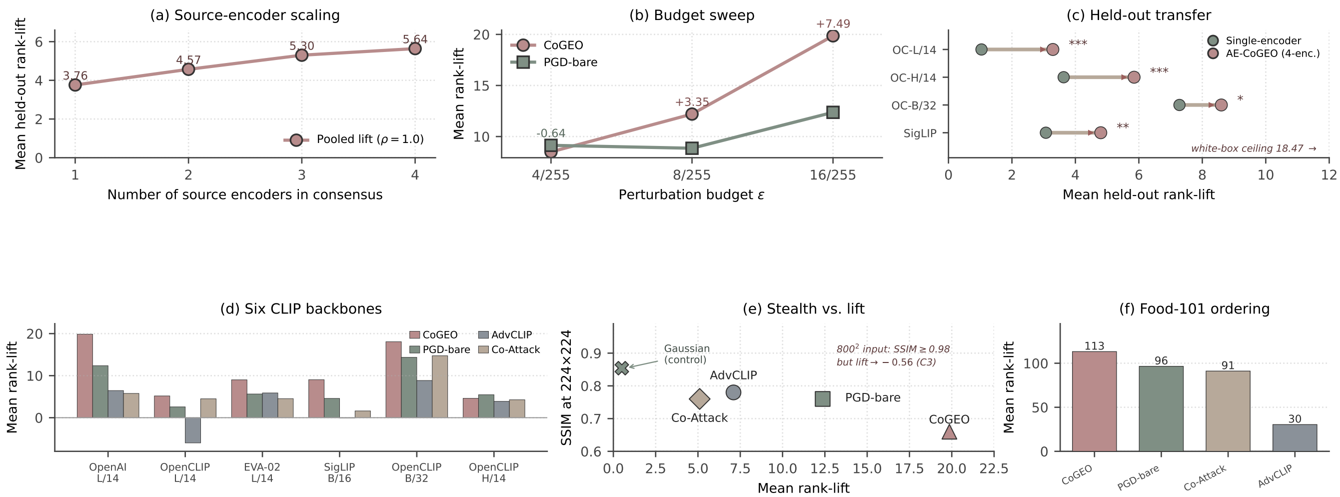

Fig. 2. AE-CoGEO transfer behavior and protocol robustness; every value matches a manuscript table. **(a)** Rank-lift scales with the source-encoder count ($3.76 \rightarrow 5.64$, $\rho = 1.0$). **(b)** Budget sweep (OpenAI ViT-L/14): the CoGEO−PGD-bare margin inverts at $\varepsilon = 4/255$ (-0.64) and grows to $+7.49$ at $16/255$. **(c)** Held-out transfer (single-encoder → four-encoder consensus, $n = 498$, Table IV): the consensus lifts all four unseen encoders (one-sided paired Wilcoxon; $^*p < 5\times10^{-2}$, $^{**}p < 10^{-2}$, $^{***}p < 10^{-3}$) but stays below the white-box ceiling (dashed). **(d)** Four-way rank-lift across six CLIP backbones ($\varepsilon = 16/255$): CoGEO is the unique top-1 on five of six. **(e)** Stealth/lift trade-off (Table II): higher rank-lift costs SSIM, all below high fidelity at 224×224 (CoGEO 0.66, AdvCLIP 0.78). **(f)** Food-101 cross-dataset check (n=2,020, OpenAI ViT-L/14): the ESCI ordering CoGEO > PGD-bare > Co-Attack > AdvCLIP is preserved (113.2/96.5/91.1/30.4; gap $1.61\times \rightarrow 1.17\times$).

TABLE III
PER-COHORT MEAN RANK-LIFT OVER THE FOUR ESCI RELEVANCE CLASSES (560-PAIR, OPENAI VIT-L/14, $\varepsilon = 16/255$): THE PER-COHORT HETEROGENEITY CLAIM (**C2**). BOLD MARKS THE PER-COHORT BEST; SIZES $n = 125$ FOR E/S/I AND 185 FOR C; VALUES AGGREGATE TO TABLE II.

Method	Exact (n=125)	Substitute (n=125)	Complement (n=185)	Irrelevant (n=125)
CoGEO	3.90	5.24	**36.77**	**25.43**
PGD-bare	**7.12**	1.84	18.86	18.53
AdvCLIP	1.00	0.94	14.46	5.53
Co-Attack	6.18	**5.47**	−1.86	16.94

A. Per-cohort rank-lift decomposition

Table III splits the same rank-lift quantity across the four ESCI cohorts; the split materially reshapes the headline. Three patterns emerge. First, CoGEO dominates where manipulation is hardest: Complement is where its machinery yields its largest gain (+36.8, $1.95\times$ PGD-bare; a small-sample peak resolving to strict $E < S < C < I$ at $3\times$ scale, §VI-I), and it also leads Irrelevant (+25.4). Second, this dominance is not uniform: on Exact, PGD-bare's mean exceeds CoGEO's by $\sim$ 3.2 positions, yet matched-pair tests find no systematic per-pair advantage (§VI-E). Third, Co-Attack alone turns negative on Complement (-1.86), its alternating branch degrading ranking when image–text correlation is already weak. The aggregate ranking thus conceals a method-specific regime: *visual GEO is a cohort-specific operator* whose lift concentrates on the harder, lower-relevance cohorts while gaining no per-pair edge over PGD-bare on Exact.

B. Pairwise paired Wilcoxon matrix

The full 4×4 paired Wilcoxon matrix (computed on all 560 paired indices pooled, not within-cohort; released artifact) has every pairwise comparison significant at $p < 10^{-3}$ except one of direction, not significance. AdvCLIP-versus-Co-Attack is that disagreement: AdvCLIP's higher mean (6.45 vs. 5.77) reverses under the one-sided paired test (Co-Attack over AdvCLIP, p=4.8×10^{-7}), so we adopt the rank-based reading. PGD-bare is significantly above both ($p < 10^{-5}$): an eval-independent anchor alone, without CoGEO's engineering, beats both better-engineered literature attacks.

C. SSIM versus rank-lift Pareto

The four SSIM means in Table II rank broadly inverse to rank-lift: the methods occupy distinct points on a roughly monotonic stealth–lift curve at $\varepsilon = 16/255$, all at low SSIM at the native 224×224 input (CoGEO 0.66, stealthiest AdvCLIP 0.78; Fig. 2e). CoGEO's SSIM ≥ 0.98 holds only at 800×800, not the regime the retriever consumes, so "high stealth and high effect" dissolves under multi-method evaluation — a trade-off structural to the budget–effect frontier.

D. Perturbation-budget sensitivity

Our main comparison fixes $\varepsilon = 16/255$, but a defender cares most about smaller budgets. Sweeping $\varepsilon \in \{4, 8, 16\}/255$ (OpenAI ViT-L/14) adds a regime-dependent qualification: at $\varepsilon = 4/255$ the headline inverts and CoGEO trails PGD-bare by 0.64, as CoGEO's auxiliary regularizers [14], [22], [23] consume a fixed share of too small a budget; as ε doubles the margin swings positive and grows (+3.35 at $8/255$, +7.49 at $16/255$; Fig. 2b). The trade-off thus shifts *within* CoGEO with the budget, so the $\varepsilon = 16/255$ ordering does not extrapolate to a smaller-budget adversary.

E. Failure mode on the Exact cohort

The Exact-cohort column of Table III is striking: PGD-bare $+7.12$, Co-Attack $+6.18$, CoGEO $+3.90$, AdvCLIP $+1.00$ — two attacks that strip CoGEO's mask and environment simulation finish above it. The matched-pair tests qualify this aggregate ordering.

Of the n=125 Exact pairs, 90 tie at rank-lift 0; among the 35 non-tie pairs 17 favor PGD-bare and 18 CoGEO (two-sided sign-test p=1.00 and paired Wilcoxon p=0.47; median paired difference 0, 10%-trimmed mean 0.01 — two orders below the raw 3.22 gap). The aggregate gap is thus not a per-pair effect but a mechanism mismatch, consistent with the Sobel-edge mask concentrating budget on high-frequency boundaries that whole-image Exact pairs do not localize. So rather than "PGD-bare beats CoGEO on Exact," the precise statement is that no method has a significant per-pair Exact advantage: CoGEO's edge is bounded to the weak image–text-alignment cohorts (Complement, Irrelevant) and statistically matched where alignment is already strong.

F. Cross-backbone replication

To check that the per-cohort dominance is an attack-family property rather than one embedding's, we replicate the four-way comparison on five more CLIP-style retrievers (§IV-B) spanning two corpora and three scales, holding manifest, queries, budget, iterations, and anchor fixed (Fig. 2d). This sharpens **(C1)** into cross-backbone form: *CoGEO is the unique top-*1 *in mean rank-lift on five of six CLIP backbones at* $\varepsilon = 16/255$*, overtaken by PGD-bare only on the highest-capacity OpenCLIP-H/14*, consistent with the Exact-cohort failure mode (§VI-E). The backbone axis is also a defensive surface: within the corpus-controlled LAION-2B family the manipulable band contracts monotonically with capacity (n=500, $\rho = -1.0$; Fig. 3); the all-six correlation is a non-significant $\rho = -0.64$.

External validity beyond ESCI/OpenAI: Food-101 is a second image-domain four-way replication (n=2,020, all four cohorts; Fig. 2f): the ranking is preserved, so the ceiling and ordering are attack-family properties, not ESCI/OpenAI artifacts. Here CoGEO's edge over PGD-bare *is* paired-significant (113.2 vs. 96.5, two-sided Wilcoxon p=1.4×10^{-7}), unlike the one-sided ESCI scale-up null below (p=0.78): the per-image-attack advantage is domain-dependent. Its cohorts are CLIP-similarity quartiles, not human grades, so it acts as a retriever-consistency probe: it corroborates ordering but cannot rule out CLIP-geometry overfitting, which the cross-architecture AE-CoGEO probe (ρ=1.0) tests independently.

G. The transfer axis: significant but bounded held-out rank-lift

AE-CoGEO is a query-free, anchor-decoupled cross-encoder probe that *measures* this transfer axis as a calibrated instrument, not a SOTA attack: one perturbation optimized against the four-encoder source consensus is scored, without re-attack, on four held-out encoders, against a single-encoder CoGEO arm (OpenAI ViT-L/14 only). The arms differ only in the source set ($n = 498$, $\varepsilon = 16/255$), isolating the transfer mechanism (§III-D).

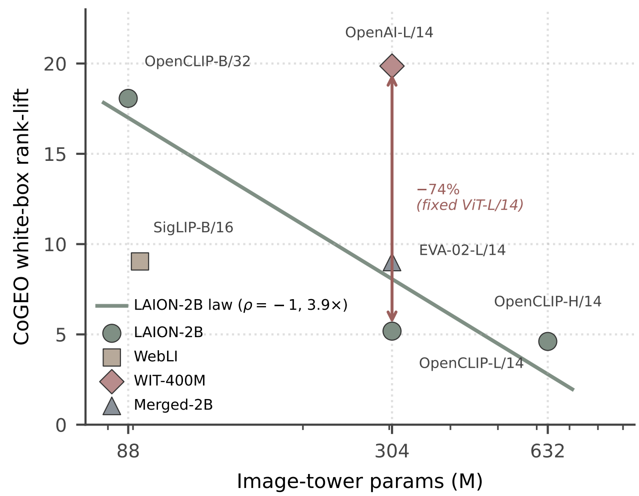

Fig. 3. Retriever capacity compresses the manipulable band. CoGEO white-box mean rank-lift versus image-tower parameter count (log scale, $\varepsilon = 16/255$, 224×224): within the corpus-controlled LAION-2B family the band contracts monotonically with capacity ($18.07 \to 5.18 \to 4.61$, Spearman $\rho = -1.0$, $\approx 3.9\times$); a larger or cleaner pretraining corpus compresses it further at fixed capacity (-74% from OpenAI WIT-400M to LAION-2B at ViT-L/14). The all-six correlation is a non-significant $\rho = -0.64$.

TABLE IV
HELD-OUT-ENCODER TRANSFER RANK-LIFT (MEAN POSITIONS GAINED, $n = 498$, $\varepsilon = 16/255$). THE SINGLE-ENCODER ARM TRANSFERS A COGEO PERTURBATION OPTIMIZED AGAINST OPENAI VIT-L/14 ALONE; THE AE-COGEO ARM TRANSFERS ONE OPTIMIZED AGAINST THE FOUR-ENCODER SOURCE CONSENSUS (§III-D); NO HELD-OUT ENCODER APPEARS IN THE SOURCE SET. p: ONE-SIDED PAIRED WILCOXON OF AE-COGEO OVER THE SINGLE-ENCODER BASELINE. †LAION-2B-PRETRAINED.

Held-out encoder	Single-enc	AE-CoGEO	p
OpenCLIP-L/14†	1.04	**3.29**	7×10^{-5}
OpenCLIP-H/14†	3.63	**5.85**	$< 10^{-3}$
OpenCLIP-B/32†	7.28	**8.60**	0.016
SigLIP ViT-B/16	3.07	**4.80**	1.3×10^{-3}
Strongest-source white-box: **18.47** (transfer ceiling)			

Consensus transfer is real and statistically significant: The consensus lifts mean held-out rank-lift over the single-encoder baseline on all four unseen encoders, each significant (one-sided paired Wilcoxon, $p \in [7 \times 10^{-5}, 0.016]$; all four clear Holm–Bonferroni at the four-test count; Table IV, Fig. 2c): an attacker ensembling public CLIP encoders measurably moves a retriever it never optimized against.

The transfer threat stays far below the white-box ceiling, across paradigms: The same perturbation evaluated white-box on its strongest source encoder reaches the transfer ceiling of 18.47 (Table IV), two-to-six times the transferred

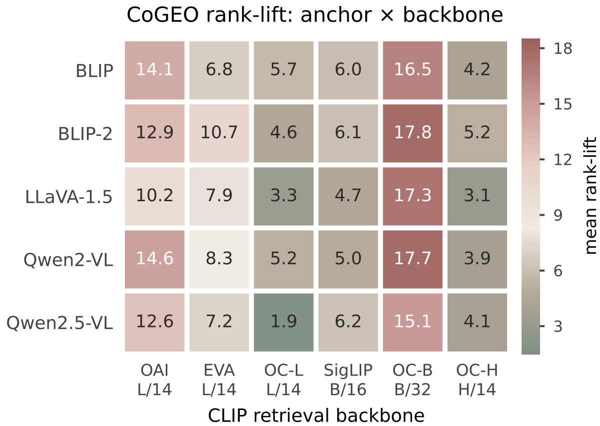

Fig. 4. CoGEO mean rank-lift across 5 VLM caption anchors (rows) $\times$ 6 CLIP backbones (columns) on the 491-ASIN ESCI manifest ($\varepsilon = 16/255$): OAI L/14, EVA (EVA-02 L/14), OC-L (OpenCLIP L/14), SigLIP (B/16), OC-B (OpenCLIP B/32), OC-H (OpenCLIP H/14). Column means span 4.1–16.9 but row means only 7.7–9.6: the retriever, not the anchor captioner, explains the bulk of the spread—the signature of an anchor-decoupled protocol.

3.3–8.6. Scored without re-attack on a fused-encoder BLIP reranker [16] sharing no encoder or objective with the source set ($n = 500$), it reaches only 7.52, inside the held-out band (Fig. 5a): the deflation survives the strongest transfer probe *and* a change of retrieval paradigm, and the detector flags it (AUROC 0.990).

Consensus transfer scales monotonically with the source-encoder count: Growing the source set from one to four public CLIP encoders raises the pooled held-out transfer monotonically ($3.76 \rightarrow 5.64$, Spearman $\rho = 1.0$; Fig. 2a), so AE-CoGEO is a genuine ensemble-transfer instrument with a measurable scaling law.

H. Anchor-source robustness

Because the protocol requires only that the anchor be a function of x_p, we re-run CoGEO against the same six retrievers with five VLM captioners, giving the 5×6 rank-lift matrix of Fig. 4. Two facts hold: the retriever dominates the anchor (per-backbone column means span 4.1–16.9, versus only 7.7–9.6 across captioners), and no captioner is uniformly best (within-column gap under 5 positions). Both confirm (1): the protocol surfaces an encoder-side property, not a captioner artifact.

I. Bounded under stress: transfer, defense, and scale-up

The remaining robustness, transfer, and defense experiments are consolidated into one display, Figure 5. Three messages emerge.

The white-box rank-lift is a genuine ceiling: Every out-of-source transfer probe — a four-encoder AE-CoGEO consensus, a BLIP fused-encoder reranker, a BLIP-2 Q-Former, and a late-interaction MaxSim matcher — stays far below it (Fig. 5a), so the deflation survives a change of encoder, of retrieval paradigm, and of matching mechanism. A fifth, diffusion-prior family (a latent-diffusion anchor-aligned perturbation; setup in the artifact) lands mid-pack (13.18, below CoGEO's 15.60 on the matched n=500 panel; Table V), so CoGEO holds the top point estimate across five families, though their CIs overlap and ranks 2–4 are sample-dependent; all five are strongly right-skewed (Fig. 6).

The bounded operator is defender-containable, and heterogeneity is not an exploitable handle: A training-free Laplacian detector flags all four families and the AE-CoGEO transfer perturbation at AUROC ≥ 0.99, generalizing across families and to a held-out domain, degrading only as the budget shrinks (Table V); input purification removes most of the residual lift, and even a matched adaptive attacker stays far below the ceiling (Fig. 5b). A cohort-adaptive-budget attacker redistributing a matched mean budget toward cheaper-to-move cohorts *under-performs* a uniform budget (13.0 vs 15.6), so heterogeneity is a diagnostic, not a handle; no tested budget is both potent and detector-evasive. Detector and purifier thus compose into a deployable audit-then-mitigate pipeline: gating JPEG purification on the detector (clean-only, 5% FPR) cuts the ε=16/255 CoGEO rank-lift $15.60 \rightarrow 7.40$ by touching only the $\sim 5\%$ flagged listings, at $20\times$ lower clean cost than blanket purification (0.51 vs. 10.07; blanket suppresses to 5.78 but at clean cost as large as the attack). At the deployable gated point $\sim 48\%$ of the lift survives (7.40 of 15.60): lightweight defenses contain but do not neutralize it — consistent with our recalibration thesis, not an easily-defeated attack.

The findings are not small-sample artifacts: Rebuilding a label-balanced ESCI sample three times larger (1,430 pairs) sharpens the per-cohort gradient and the mean-shift, not per-pair directional dominance (Fig. 5c). The mean ordering is CoGEO 41.26 > Co-Attack 34.05 ≈ PGD-bare 33.38 > AdvCLIP 26.16. The CoGEO−PGD-bare lead is +7.89 (bootstrap 95% CI $[2.0, 13.6]$) yet the paired directional Wilcoxon is non-significant (p=0.78; over Co-Attack, p=0.67), so the advantage over both is tail-driven — though CoGEO does paired-significantly exceed the prior published AdvCLIP attack at this scale ($p < 10^{-40}$, every cohort). Both anchor attacks rise strictly ($E < S < C < I$), resolving the small-sample 560-pair Complement-peak ordering (Table III). Two directional controls reinforce this: Gaussian noise at the same budget is stealthier than every attack (SSIM 0.854) yet yields ≈ 0 lift, and the same attack at 800^2 input is more stealthy (SSIM ≥ 0.98) but loses essentially all lift, inverting to -0.56 versus PGD-bare's $+1.22$ (C3; Fig. 2e).

VII. Limitations

The empirical scope is intentionally narrow. **Architecture.** Our six retrievers [1], [10]–[12] span two pretraining corpora and three scales but share the CLIP dual-encoder paradigm; §VI-G measures bounded positive transfer to a BLIP fused-encoder reranker [16] (7.52, within the 3.3–8.6 band), no positive transfer to a BLIP-2 Q-Former [30], and the same bounded behavior on a MaxSim probe; trained late-interaction and learned-hashing retrievers are future work. The encoder pair is both optimization target and evaluation embedder, so headline numbers are a white-box upper bound. **Resolution and budget.** Headline rankings use 224×224 at $\varepsilon = 16/255$;

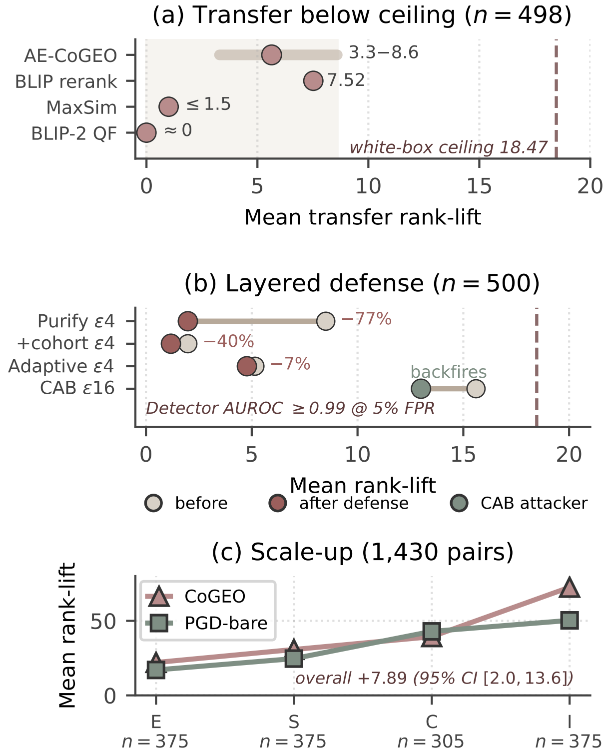

Fig. 5. Consolidated robustness, transfer, and defense evidence (rank-lift vs. the 18.47 white-box ceiling; panel (c) is an independent 1,430-pair scale-up with its own axis). **(a)** Out-of-source transfer: the AE-CoGEO four-encoder consensus (3.3–8.6), a BLIP reranker (7.52), a MaxSim probe (≤ 1.5), and a BLIP-2 Q-Former (≈ 0) all sit far below the ceiling. **(b)** Layered defense at ε=4/255 (before $\rightarrow$ after): input purification (8.50 $\rightarrow$ 1.98), cohort-aware oracle (-40% further), a matched adaptive attacker (5.15 $\rightarrow$ 4.77); at ε=16/255 a cohort-adaptive-budget attacker under-performs uniform (13.0 vs 15.6). **(c)** The 3× scale-up sharpens the per-cohort gradient into a strictly monotone $E < S < C < I$ for both load-bearing attacks (CoGEO−PGD +7.89, 95% CI $[2.0, 13.6]$). Values match §VI-I.

the artifact reports two crossovers (800×800 and $\varepsilon = 4/255$, both favoring PGD-bare), so CoGEO's lead requires a budget large enough to amortize its auxiliary objectives. **Attack and benchmark scope.** The four-attack shortlist [4], [5], [8], [18] covers the per-image gradient, universal-patch, and cross-modal families. A gradient-free NES black-box probe (2,000 forward queries/image, anchor-decoupled) reaches only 3.05 held-out rank-lift versus CoGEO's 13.4 on a matched 120-pair subset under identical scoring (23%; two-sided paired Wilcoxon $p < 10^{-3}$), so the white-box lead is a genuine attack-strength bound, not a gradient-access artifact. Consistent with the per-cohort picture, CoGEO's edge is bounded to weak image–text-alignment cohorts and is statistically matched on Exact — a cohort-specific operator, not a uniform one. Crucially, the human-relevance findings rest on ESCI's graded labels [7] alone (the Food-101 cohorts are retriever-induced, §VI-F): no surveyed public benchmark offers a second product domain with multi-level human relevance and usable images. **Anchor independence.** The invariant of (1) closes the immediate optimization-anchor/eval-query loop but not the deeper seller-ecosystem/training-corpus loop; the

TABLE V

TWO FURTHER CHECKS (OPENAI VIT-L/14, n=500). *Top*: PER-PAIR RANK-LIFT FOR FIVE WHITE-BOX FAMILIES AT ε=16/255; ALL ARE RIGHT-SKEWED (MEDIAN $\approx$ 1) WITH COGEO THE TOP POINT ESTIMATE (OVERLAPPING CIS). *Bottom*: TRAINING-FREE LAPLACIAN-DETECTOR AUROC BY BUDGET, $\geq$ 0.99 AT ε=16/255 AND DEGRADING ONLY AS THE BUDGET SHRINKS.

Five attack families (white-box), per-pair rank-lift					
Family	Mean	95% CI	Med.	$\% \leq 0$	Max
CoGEO	**15.60**	[10.9, 20.3]	1.0	47.0	354
PGD-bare	13.58	[8.3, 18.8]	1.0	48.2	390
Diffusion-prior	13.18	[8.5, 17.9]	1.0	48.8	388
Co-Attack	12.82	[8.2, 17.5]	1.0	48.0	348
AdvCLIP	5.06	[1.8, 8.4]	0.0	63.4	281
Training-free detector AUROC by budget					
ε=16/255	0.99–1.00		ε=8/255: 0.957–0.970		
ε=4/255	0.852–0.889		Food-101: 1.00		

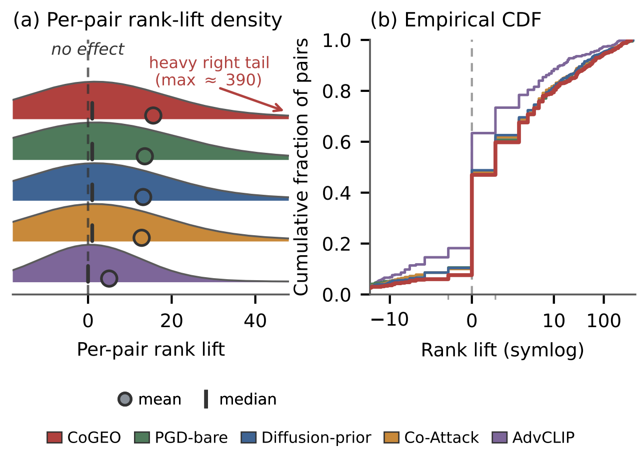

Fig. 6. Per-pair rank-lift distribution for the five white-box attack families (OpenAI ViT-L/14, n=500 each, ε=16/255). **(a)** Ridgeline densities: every family is strongly right-skewed—the typical pair barely moves (median tick at $\approx$1; dashed "no effect" line) while a heavy right tail (up to $\approx$390) lifts the mean (circle markers; shared key below). **(b)** Empirical CDF on a symlog axis (family colours as in (a)): the lower/further-right a curve, the heavier its tail; CoGEO is heaviest, AdvCLIP lightest, with $\geq 47\%$ of pairs unaffected (≤ 0) for every family. The headline mean thus summarises a sparse, skewed effect, not a uniform shift; the marginals overlap, but the headline 560-pair test is on *paired* per-pair deltas (Wilcoxon $p \leq 1.1 \times 10^{-4}$), not on these marginals.

VLM-anchor ablation (released artifact) bounds within-anchor variation by under 5 rank positions across five captioners and six backbones, narrowing the self-loop critique rather than erasing the per-cohort findings.

VIII. CONCLUSION

Within the anchor-decoupled protocol — an anchor$\perp$query invariant, a pre-flight retriever-discrimination gate, and a per-cohort decomposition over the four ESCI classes [7] — a four-way comparison on Amazon ESCI across six retrievers at $\varepsilon = 16/255$ replaces optimistic self-loop scores with honest measurement. Our CoGEO is the strongest of the four attacks: the unique top-1 on five of six retrievers, dominant on the harder Complement (+36.8) and Irrelevant (+25.4) cohorts, and paired-significantly ahead of plain PGD on the Food-101 cross-domain replication (113.2 vs. 96.5). The per-cohort lens shows this strength is concentrated where image–text alignment is weak rather than spread uniformly, turning a single inflated headline into a calibrated, defender-actionable threat map.

Beyond the white-box axis, AE-CoGEO — the first cross-encoder transfer probe for visual GEO — lifts rank on all four held-out encoders with a clean monotone source-count scaling law (Spearman $\rho = 1.0$), calibrating what a deployed retriever actually sees. The same instruments contain the threat: a training-free detector flags every attack family at $\mathrm{AUROC} \geq 0.99$ and detector-gated purification removes most of the residual lift, composing into an *audit-then-mitigate* defense. Visual GEO thus reads as a bounded, cohort-specific, detectable operator, not a uniform invisible threat. We release the protocol and both attacks anonymously as HONESTEVAL (https://anonymous.4open.science/r/honestvgeo-F663), evaluated only on public ESCI data; the open gap is a second human-graded relevance domain beyond our retriever-induced Food-101 cohorts.

REFERENCES

[1] A. Radford, J. W. Kim, C. Hallacy, A. Ramesh, G. Goh, S. Agarwal, G. Sastry, A. Askell, P. Mishkin, J. Clark, G. Krueger, and I. Sutskever, "Learning transferable visual models from natural language supervision," in *Proceedings of the 38th International Conference on Machine Learning (ICML)*, 2021, pp. 8748–8763.

[2] P. Aggarwal, V. Murahari, T. Rajpurohit, A. Kalyan, K. Narasimhan, and A. Deshpande, "GEO: Generative engine optimization," in *Proceedings of the 30th ACM SIGKDD Conference on Knowledge Discovery and Data Mining (KDD)*, 2024.

[3] I. J. Goodfellow, J. Shlens, and C. Szegedy, "Explaining and harnessing adversarial examples," in *International Conference on Learning Representations (ICLR)*, 2015.

[4] Z. Zhou, S. Hu, M. Li, H. Zhang, Y. Zhang, and H. Jin, "AdvCLIP: Downstream-agnostic adversarial examples in multimodal contrastive learning," in *Proceedings of the 31st ACM International Conference on Multimedia*, 2023, pp. 6311–6320.

[5] J. Zhang, Q. Yi, and J. Sang, "Towards adversarial attack on vision-language pre-training models," in *Proceedings of the 30th ACM International Conference on Multimedia*, 2022, pp. 5005–5013.

[6] Z. Wang, A. C. Bovik, H. R. Sheikh, and E. P. Simoncelli, "Image quality assessment: From error visibility to structural similarity," *IEEE Transactions on Image Processing*, vol. 13, no. 4, pp. 600–612, 2004.

[7] C. K. Reddy, L. Màrquez, F. Valero, N. Rao, H. Zaragoza, S. Bandyopadhyay, A. Biswas, A. Xing, and K. Subbian, "Shopping queries dataset: A large-scale ESCI benchmark for improving product search," *arXiv preprint arXiv:2206.06588*, 2022.

[8] A. Madry, A. Makelov, L. Schmidt, D. Tsipras, and A. Vladu, "Towards deep learning models resistant to adversarial attacks," in *International Conference on Learning Representations (ICLR)*, 2018.

[9] F. Wilcoxon, "Individual comparisons by ranking methods," *Biometrics Bulletin*, vol. 1, no. 6, pp. 80–83, 1945.

[10] M. Cherti, R. Beaumont, R. Wightman, M. Wortsman, G. Ilharco, C. Gordon, C. Schuhmann, L. Schmidt, and J. Jitsev, "Reproducible scaling laws for contrastive language-image learning," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.

[11] Y. Fang, Q. Sun, X. Wang, T. Huang, X. Wang, and Y. Cao, "EVA-02: A visual representation for neon genesis," *arXiv preprint arXiv:2303.11331*, 2023.

[12] X. Zhai, B. Mustafa, A. Kolesnikov, and L. Beyer, "Sigmoid loss for language image pre-training," in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023.

[13] C. Jia, Y. Yang, Y. Xia, Y.-T. Chen, Z. Parekh, H. Pham, Q. V. Le, Y.-H. Sung, Z. Li, and T. Duerig, "Scaling up visual and vision-language representation learning with noisy text supervision," in *Proceedings of the 38th International Conference on Machine Learning (ICML)*, 2021, pp. 4904–4916.

[14] R. Shin and D. Song, "JPEG-resistant adversarial images," in *NeurIPS Workshop on Machine Learning and Computer Security*, 2017.

[15] Y. Dong, F. Liao, T. Pang, H. Su, J. Zhu, X. Hu, and J. Li, "Boosting adversarial attacks with momentum," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2018, pp. 9185–9193.

[16] J. Li, D. Li, C. Xiong, and S. C. H. Hoi, "BLIP: Bootstrapping language-image pre-training for unified vision-language understanding and generation," in *Proceedings of the 39th International Conference on Machine Learning (ICML)*, 2022.

[17] P. Wang, S. Bai, S. Tan, S. Wang, Z. Fan, J. Bai, K. Chen, X. Liu, J. Wang, W. Ge, Y. Fan, K. Dang, M. Du, X. Ren, R. Men, D. Liu, C. Zhou, J. Zhou, and J. Lin, "Qwen2-VL: Enhancing vision-language model's perception of the world at any resolution," *arXiv preprint arXiv:2409.12191*, 2024.

[18] S.-M. Moosavi-Dezfooli, A. Fawzi, O. Fawzi, and P. Frossard, "Universal adversarial perturbations," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2017, pp. 1765–1773.

[19] A. Ilyas, S. Santurkar, D. Tsipras, L. Engstrom, B. Tran, and A. Madry, "Adversarial examples are not bugs, they are features," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2019.

[20] N. Carlini and D. A. Wagner, "Towards evaluating the robustness of neural networks," in *IEEE Symposium on Security and Privacy*, 2017, pp. 39–57.

[21] C. Schlarmann, N. D. Singh, F. Croce, and M. Hein, "Robust CLIP: Unsupervised adversarial fine-tuning of vision embeddings for robust large vision-language models," in *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024, arXiv:2402.12336.

[22] C. Xie, Z. Zhang, Y. Zhou, S. Bai, J. Wang, Z. Ren, and A. L. Yuille, "Improving transferability of adversarial examples with input diversity," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019, pp. 2730–2739.

[23] J. Lin, C. Song, K. He, L. Wang, and J. E. Hopcroft, "Nesterov accelerated gradient and scale invariance for adversarial attacks," in *International Conference on Learning Representations (ICLR)*, 2020.

[24] N. Papernot, P. McDaniel, I. Goodfellow, S. Jha, Z. B. Celik, and A. Swami, "Practical black-box attacks against machine learning," in *Proceedings of the ACM Asia Conference on Computer and Communications Security*, 2017, pp. 506–519.

[25] L. Bossard, M. Guillaumin, and L. Van Gool, "Food-101 – mining discriminative components with random forests," in *Proceedings of the European Conference on Computer Vision (ECCV)*, 2014.

[26] N. Thakur, N. Reimers, A. Rücklé, A. Srivastava, and I. Gurevych, "BEIR: A heterogeneous benchmark for zero-shot evaluation of information retrieval models," in *Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2021.

[27] N. Muennighoff, N. Tazi, L. Magne, and N. Reimers, "MTEB: Massive text embedding benchmark," in *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics (EACL)*, 2023, pp. 2014–2037.

[28] C. Mao, S. Geng, J. Yang, X. Wang, and C. Vondrick, "Understanding zero-shot adversarial robustness for large-scale models," in *The Eleventh International Conference on Learning Representations (ICLR)*, 2023.

[29] S. Thys, W. Van Ranst, and T. Goedemé, "Fooling automated surveillance cameras: Adversarial patches to attack person detection," in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR) Workshops*, 2019.

[30] J. Li, D. Li, S. Savarese, and S. Hoi, "BLIP-2: Bootstrapping language-image pre-training with frozen image encoders and large language models," in *International Conference on Machine Learning (ICML)*, 2023.

[31] H. Liu, C. Li, Y. Li, and Y. J. Lee, "Improved baselines with visual instruction tuning," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.

[32] S. Bai, K. Chen, X. Liu, J. Wang, W. Ge, S. Song, K. Dang, P. Wang, S. Wang, J. Tang *et al.*, "Qwen2.5-VL technical report," *arXiv preprint arXiv:2502.13923*, 2025.